

UNIX PROGRAMMING ASSIGNMENT

MCA - 372

Submitted by :- Riya Gurbani

Roll Number : 1947255

Describe the following from Process Control of Unix Environment

1) Fork, and Vfork with file sharing functions
Fork

When we create a ^{new} process through an existing one is by calling the fork function.

When we redirect the standard output of the parent from the program, the child's standard output is also redirected. All file descriptions that are open in the parent are duplicated in the child. Because it's as if the dup function had been called for each descriptor. The parent & the child share a file table entry for every open descriptor. Consider a process that has three different files opened for standard-input, output & error. On return from fork, we have the arrangement.

It is important that the parent & the child share the same file offset. Consider a process that forks a child, then waits for the child to complete.

Assume that both processes write to standard output as part of their normal processing.

If the parent has its standard output redirected it is essential that the parent's file offset be updated by the child when the child writes to

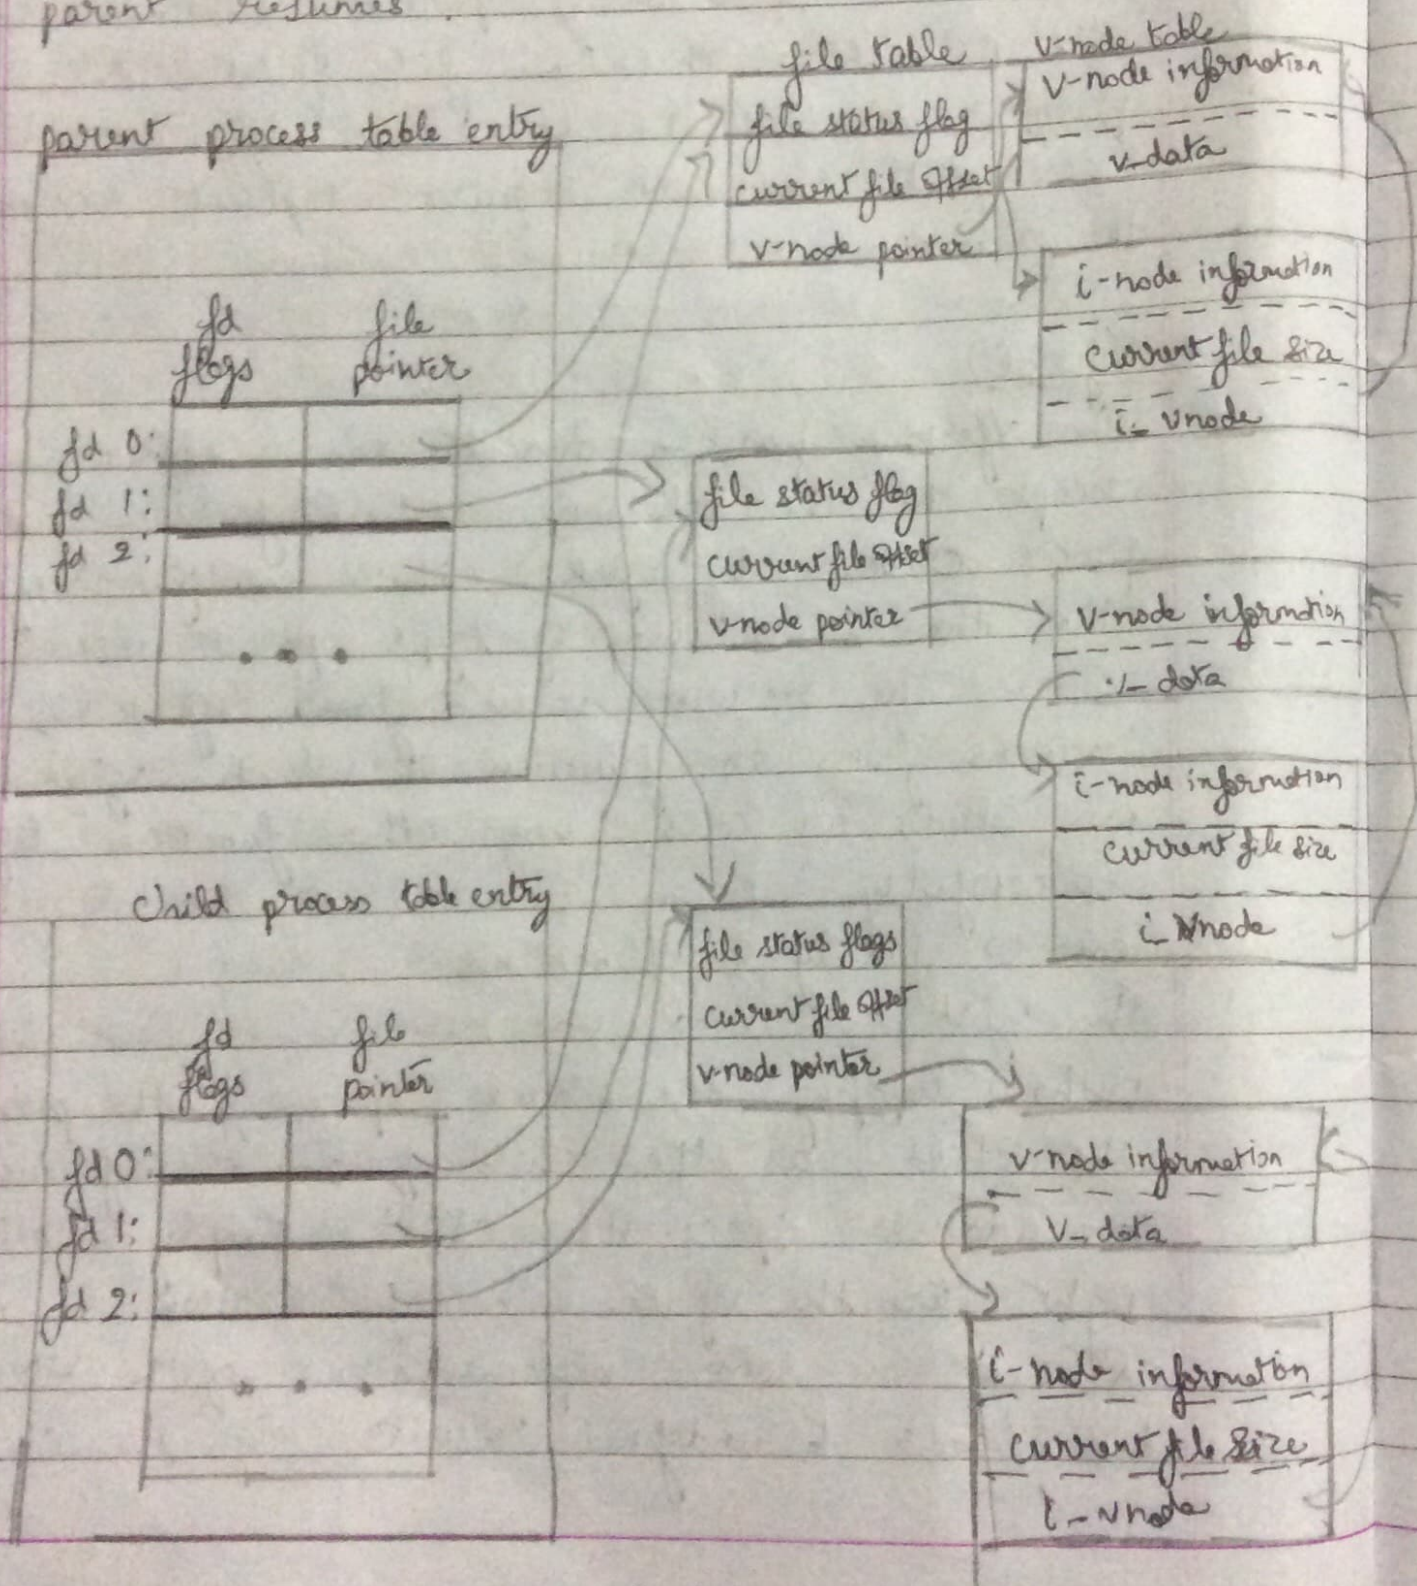
standard output. In this case, the child can write to standard output while the parent is waiting for it. on completion of the child, the parent can continue writing to standard output, knowing that its output will be appended to whatever the child wrote. If parent & the child did not share the same file offset, this type of interaction would be more difficult to accomplish & would require explicit actions by the parent.

Vfork()

The `vfork()` function has the same effect as `fork()`, except that the behaviour is undefined if the process created by `vfork()` either modifies any data other than a variable of type `pid_t` used to store the return value from `vfork()`, or returns from the function in which `vfork()` was called, or calls any other function before successfully calling `_exit()` or one of the `_exec()` family of functions.

The `vfork()` creates the new process just like `fork`, without copying the address space of the parent into the child, as the child won't reference that ~~an~~ address space, the child simply calls `_exec` right after the `vfork`. The optimization is more efficient on some implementations of the UNIX system, but leads to undefined results if the child

modifies any data, makes function calls, or return without calling `exec` or `exit`.
`Vfork` guarantees that the child runs first, until the child calls `exec` or `exit`. When the child calls either of those functions, the parent resumes.



Sharing of open files between parent & child after `fork`.

2. Wait and Wait ID functions.

When a process terminates, either normally or abnormally, the kernel notifies the parent by sending the `SIGCHLD` signal to the parent. Because the termination of a child is an asynchronous event, it can happen at any time while the parent is running - this signal is the asynchronous notification from the kernel of the parent. The parent can choose to ignore this signal, or it can provide a function that is called when the signal occurs: a signal handler. The default action for this signal is to be ignored.

- Block, if all of its children are still running
- Return immediately with the termination status of a child, if a child has ~~been~~ terminated & is waiting for its termination status to be fetched.
- Return immediately with an error, if it doesn't have any child processes.

If the process is calling wait because it received the `SIGCHLD` signal, we expect wait to return immediately. But if we call it any random point in time, it can block.

```
#include <sys/wait.h>
```

```
pid_t wait (int *statloc);
```

```
pid_t waitpid (pid_t pid, int *statloc, int options);
```

The difference b/w these two functions are:-

- The wait function can block the caller, until a child process terminates, whereas waitpid has an

option that prevents it from blocking.

→ The `waitpid()` doesn't wait for the child that terminates first, it has a number of options that control which process it waits for.

If a child has already terminated into ξ , it's a zombie, `wait` returns immediately with that child's status. Otherwise it blocks the caller until a child terminates. If the caller blocks & has multiple children, `wait` returns when one terminates. We can always tell which child terminated, because the process ID is returned by the function.

waitid function

The Single Unit Specification includes an additional function to retrieve the exit status of a process. The `waitid()` is similar to `waitpid` but with extra flexibility.

```
#include <sys/wait.h>
```

```
int waitid(idtype_t idtype, id_t id, siginfo_t  
* info, int options);
```

Return 0 if ok -1 on error

Like `waitpid`, `waitid` allows a process to specify which of its children to wait for. Instead of encoding this information in a single argument

combined with the process ID or process group ID, two separate arguments are used. The id parameter is interpreted based on the value of idtype.

The types supported are:-

Constant	Description
P_PID	wait for a particular process: id contains the process ID of the child to wait for.
P_PGID	wait for any child process in a particular process group: id contains the process group ID of the children to wait for.
P_ALL	wait for any child process: id is ignored.

info A pointer to a siginfo_t structure, where the function can store the current state of the child.

3 exec function:-

When a process calls one of the exec functions, that process is completely replaced by the new program, and the new program starts executing at its main function. The process ID does not change across an exec, because a new process is not created; exec merely replaces the current process - its text, data, heap, & stack segments - with a brand new program from disk. There are seven different exec functions, round out the UNIX system process control primitives. With fork, we can create a new process & with exec functions, we can initiate new programs.

The `exec` function & the `wait` functions handle termination & waiting for termination.

```
#include <unistd.h>
int exec (const char *pathname, const char *arg0,
          /* (char *) 0 */);
int execl (const char *pathname, char *const argv[]);
int execl (const char *pathname, const char *arg0, ---
int execl (const char *pathname, char *const argv[],
char *const envp[]);
```

In all the above examples the first four take a `pathname` argument, the next two take a `filename` argument, and the last one takes a file descriptor argument. When a `filename` argument is specified,

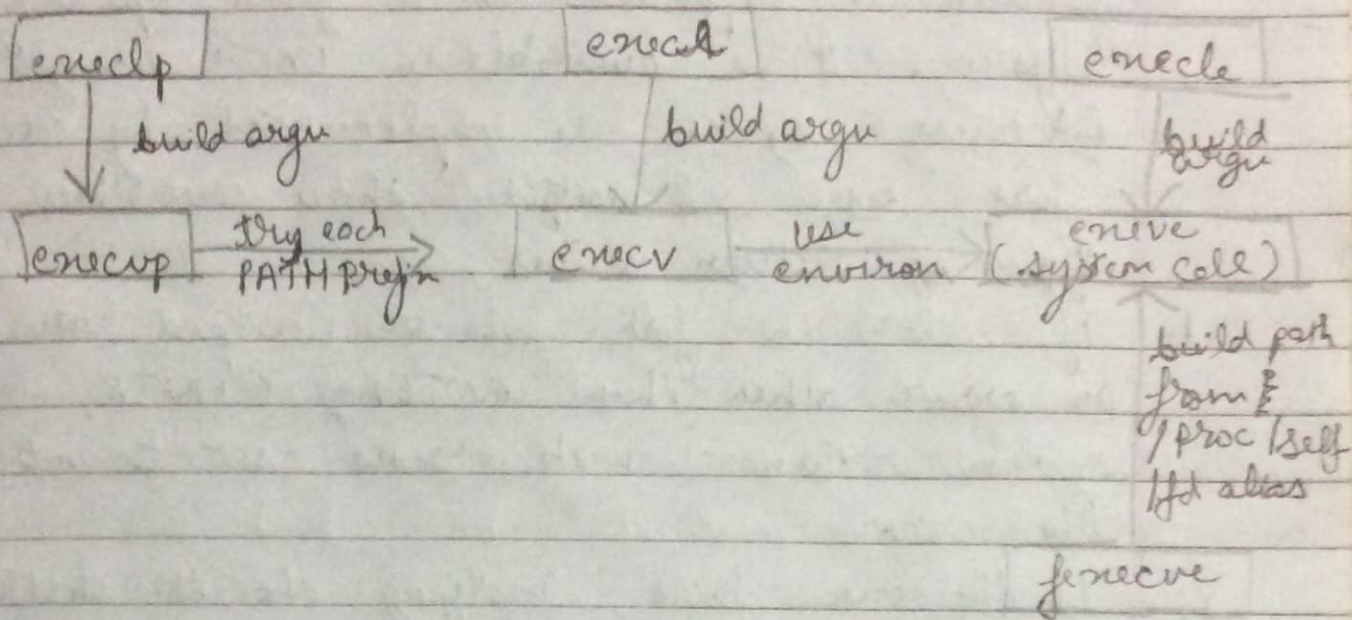
→ If `filename` contains a slash, it is taken as a `pathname`.

→ Otherwise, the executable file is searched for in the directories specified by the `PATH` environment variable.

Second is the argument passing (l stands for list & v stands for vector).

Since the arguments for these `exec` functions are difficult to remember, the letters in the function names help somehow. The letter `p` means that the function takes a `filename` argument & uses the `PATH` env. variable to find the executable file.

The letter λ means it takes a list of arguments
 λ is mutually exclusive with the letter v .
 the letter e means that the function takes an
 envp [] array instead of using the current
 environment.



Relationship of the seven exc functions.

4

System Function

It is convenient to execute a command string from within a program. For eg. if we want to display time & date stamp into a certain file we can use the function `system("date > file")`. Call `time` to get current calendar time, then call `localtime` to convert it to a broken-down time, then call `strftime` to format the result & finally write the result to the file.


```
#include <stdlib.h>
```

```
int system(const char *cmdstring);
```

If `cmdstring` is a null pointer, `system` returns nonzero only if a command processor is available. It determines whether the system function is supported on a given operating system. It is available in UNIX always.

Because `system` is implemented by calling `fork`, `exec`, & `waitpid`, there are three types of return values.

1. If either the `fork` fails or `waitpid` returns an error other than `EXITED`, `FINTR`, `system` returns -1 with errors set to indicate the error.
2. If the `exec` fails, implying that the shell can't be executed, the return value is 0 if the shell had executed `exit(127)`.
3. Otherwise all three functions - `fork`, `exec`, & `waitpid` - succeed, & the return value from `system` is the termination status of the shell, in the format specified for `waitpid`.

Set-User-ID Programs.

It creates a security hole & should never be attempted.

5. Process Scheduling

The scheduling policy and priority were determined by the kernel. A process could choose to run with lower priority by adjusting its nice value. Only a privileged process was allowed to increase its scheduling priority. In the Single Unix specification, nice values range from 0 to $(2 * NZERO) - 1$, lower nice values have higher scheduling priority. NZERO is the default nice value of the system.

A process can retrieve & change its nice value with the nice function. With this function, a process can affect only its own nice value, it can't affect the nice value of any other process.

```
#include <unistd.h>
```

```
int nice (int incr);
```

incr arg. is added to the nice value of the calling process. If incr is too large, the system silently reduces it to the minimum legal value. If it's too small the system silently increases it to the minimum legal value. Because -1 is a legal successful return value, we need to clear errors before calling nice & check its value if nice returns -1. If the call to nice succeeds & the return value is -1, then errno will still be zero. If errno is nonzero

it means that the call to nice failed.

The `getpriority()` can be used to get the nice value for a process, just like the `nice` function. However, `getpriority` can also get the nice value for a group of related processes.

```
#include <sys/resource.h>
```

```
int getpriority (int which, id_t who)
```

Returns: nice value $\in [-NZERO, NZERO-1]$
if OK, -1 on error, error.

The `which` arg can take on one of 3 values: `PRIO_PROCESS` to indicate process, `PRIO_PGID` to indicate a process group & `PRIO_USER` to indicate a user ID. The `which` arg controls how the `who` arg is interpreted & the `who` arg selects the process or processes of interest.

The `setpriority()` can be used to set the priority of a process, a process group, or all the processes belonging to a particular user ID.

```
#include <sys/resource.h>
```

```
int setpriority (int which, id_t who, int value)
```

Returns: 0 if OK, -1 on error

The `which` & `who` arg are the same as in the `getpriority()`. The `value` is added to `NZERO`.

∴ this becomes the new nice value.

The Single UNIX specification leaves it up to the implementation whether the nice value is inherited by a child process after a fork. However, XSI-compliant systems are required to preserve the nice value across a call to `exec`.