

Describe the following from Process Control of Unix Environment

1) Fork and Vfork with file sharing functions

Fork:-

When we create a new process through an existing one is by calling the fork function.

When we redirect the standard output of the parent from the program, the child's standard output is also redirected.

All file descriptions that are open in the parent are duplicated in the child, because it's as if the duplicate function had been called for each descriptor. The parent and the child share a file table entry for every open descriptor. Consider a process that has three different files opened for standard input, output & error. On return from fork, we have the arrangement

It is important that the parent & the child share the same file offset. Consider a process that forks a child, then waits for the child to complete. Assume that both processes write to standard output as part of their normal processing. If the parent has its standard output redirected it is essential that the parent's file be updated by the ~~child~~ child when the child writes to standard output.

In this case the child can write to standard output

(2)

While the parent is waiting for its ~~complication~~ completion of the child, the parent can continue writing to standard output, knowing that its output will be appended to whatever the child wrote. If parent & the child did not share the same file offset, the type of interaction would be more difficult to accomplish & would require explicit actions by the parent.

`vfork()`

The `vfork()` function has the same effect as `fork()` except that the behaviour is undefined if the process created by `vfork()` either modifies any data other than a variable of type `pid_t` used to store the return value from `vfork()` or returns from the function in which `vfork()` was called, or calls any other function before successfully calling `exit()` or one of the `exec()` family of functions.

The `vfork()` creates the new process just like `fork()` without copying the address space of the parent into the child, as the child won't reference that address space the child simply calls `exec` right after `vfork`. The optimization is more efficient on some implementations of the UNIX system, but leads to undefined results if the child modifies any data, makes function calls or returns without calling `exec` or `exit`.

`vfork` guarantees that the child runs first until the child calls `exec` or `exit`. When the child calls either of those functions the parent resumes.

2) Wait and WaitID functions

When a process terminates, either normally or abnormally, the kernel notifies the parent by sending the SIGCHLD signal to the ~~parent~~ parent, because the termination of a child is an asynchronous event, it can happen at any time while the parent is running this signal is the asynchronous notification from the kernel of the parent. The parent can choose to ignore the signal, or it can provide a function that is called when the signal occurs a signal handler. The default action for this signal is to be ignored.

- * Block if its children are still running.
- * Return immediately with the termination status of a child, if a child has terminated & is waiting for its termination status to be fetched.
- * Return immediately with an error, if it doesn't have any child processes.

If the process is calling wait because it received the SIGCHLD signal, we expect wait to return immediately but if we call it any random point in time, it can block.

```
#include <sys/wait.h>
```

```
pid_t wait (int *statloc);
```

```
pid_t waitpid (pid_t pid, int *statloc, int options);
```

The difference between these two functions are:-

- * The wait function can block the caller until a child process terminates, whereas waitpid has an option that prevents it from blocking.

* The `waitpid()` doesn't wait for the child that terminates first, it has a number of options that control which process it waits for.

If a child already terminated & a zombie, `wait` returns immediately with that child's status. Otherwise it blocks the caller ~~for~~ until a child terminates. If the caller blocks and has multiple children, `wait` returns when one terminates. We can always tell which child terminated because the process ID is returned by the function.

Waitid function

The single unit specification includes an additional function to retrieve the exit status of a process. The `waitid()` is similar to `waitpid` but with extra flexibility.

```
#include <sys/wait.h>
```

```
int waitid(idtype_t idtype, id_t id, siginfo_t * info, int options);
```

Return 0 if OK -1 on error.

Like `waitpid`, `waitid` allows a process to specify which children to wait for. Instead of encoding this information in a single argument combined with the process ID or process group ID, two separate arguments are used. The `id` parameter is interpreted based on the value of `idtype`.

The types supported are:

Constant	Description
----------	-------------

P_PID	wait for a particular process, <code>id</code> contains the process ID of the child to wait for.
-------	--

P_PGID	wait for any child process in a particular process group. <code>id</code> contains the process group ID of the children to wait for.
--------	--

P All Wait for any child process id is ignored

info A pointer to a `siginfo_t` structures where the function can store the current state of the child.

3) Exec function

When a process calls one of the exec functions that process is completely replaced by the new program and the new program starts executing at its main function. The process Id does not change across an exec, because a new ^{process} is not created. exec merely replaces the current process, its data, heap & stack segments with a brand new program from disk. There are seven different exec functions, round out the UNIX system process control primitives. With fork, we can create a new process & ^{with} exec functions we can initiate new programs.

The exit function the wait function handle termination and waiting for termination.

```
#include <unistd.h>
```

```
int execl(const char *pathname, const char *arg0,  
        ..., /* (char *) 0 */);
```

```
int exec(const char *pathname, const char *const arg[]);
```

```
int execl(const char *pathname, const char *arg0, ...
```

```
int exec(const char *pathname, char *const arg[],  
        char *const envp[]);
```

In all the above example the first four take a pathname argument the next two take a filename arg. When a filename argument is specified.

* If filename contains a slash, its taken as a pathname

* Otherwise, the executable file is searched for in the directories specified by the PATH environment variable

Second is the argument passing (l stands for list & v stands for vector)

Since the arguments for these exec functions are different to remember. The letters in the function names help

Somehow, The letter P means that the function takes a filename argument & uses the path environment variable to find the executable file.

The letter l means it takes a list of arguments and is mutually exclusive with the letter v. letter e means that the function takes an env [] array instead of using the current environment.

4) System Function

It is convenient to execute a command string from within a program. For eg: if we want to display time & date stamp into a certain file we can use the function `system('date > file')`, call `time` to get current calendar time, then call `localtime` to convert it to a broken-down time, then call `strftime` to format the result & finally write the result to the file.

```
#include <stdlib.h>
```

```
int system(const char *cmdstring);
```

If `cmdstring` is a null pointer, `system` returns non zero only if command processor is available. It determines whether the `system` function is supported on a given operating system. It is available in UNIX always because `system` is implemented by calling `fork`, `exec` & `waitpid`, there are ~~three~~ 3 types of return values.

- ① If either the `fork` fails or `waitpid` returns on error other than `EINTR`, `system` returns `-1` with error set to indicate the error.
- ② If the `exec` fails, implying that the shell can't be executed the return value is if the shell had ~~an~~ executed `exit(127)`
- ③ Otherwise all three functions - `fork`, `exec`, & `waitpid`, succeed, & the return value for from `system` is the termination status of the shell in the format specified for `waitpid`.

Set User-ID Programs

It creates a security hole & should never be attempted.

5) Process Scheduling

The scheduling policy and priority were determined by the kernel. A process could choose to run with lower priority by adjusting its nice value. Only a privileged process was allowed to increase its scheduling priority. In the ~~sig~~ Single Unix specification nice values range from 0 to $(2 \times \text{NZERO}) - 1$. Lower nice values have higher scheduling priority. 'NZERO' is the default nice value of the system.

A process can retrieve and change its nice value with the nice function. With this function a process can affect only its own nice value, it can't affect the nice value of any other process.

```
#include <unistd.h>
```

```
int nice (int incr);
```

incr arg is added to the nice value of the calling process. If incr is too large the system silently reduces it to the maximum legal value. If it's too small the system silently increases it to the minimum legal values. Because -1 is a legal successful return value, we need to clear error before calling nice & check its value if nice returns -1. If the call to nice succeeds & the return value is -1 then error will still be zero. If error is non-zero it means that the call to nice failed. The getpriority() can be used to get the nice value for a process, just like the nice function. However getpriority can also get the nice value for a group of related processes.

```
#include <sys/resource.h>
```

```
int getpriority (int which, id_t who)
```

Returns nice value between NZERO & NZERO-1

if OK, -1 on error

The which arg can take on one of 3 values PRIO_PROCESS to indicate process, PRIO_PGID to indicate a process group & PRIO_USER to indicate a user ID. The which arg controls how the who arg is interpreted and the who arg selects the process or processes of interest.

The setpriority() can be used to set the priority of a process a group process group, or all the processes belonging to a particular user Id.

```
#include <sys/resource.h>
```

```
int setpriority (int which, id_t who, int value);
```

Returns 0 if ok, -1 on error.

The which and who are the same as in the getpriority()

The value is added to NZERO and this becomes the new nice value.

The single UNIX specification leaves it up to the implementation whether the nice value is inherited by a child process after a fork. However, LSI complaints systems are required to preserve the nice value across a call to exec.